Data-Aware Networking: Routing queries in a federated environment

Jordan Augé   Marc-Olivier Buob   Loïc Baron   Timur Friedman   Serge Fdida

1. INTRODUCTION

The amount of data that is produced today in various sectors, including the scientific community, is growing at an unprecedented rate. Numerous applications require that data belonging to independent organization can be distributed over the network to be processed, combined and analyzed. The measurement and monitoring of Internet paths, crossing multiple domains, is an instance of such an issue, requiring cooperation by several independent entities. Much research so far has consisted in proposing new architecture and protocols designed for specific federations.

In this paper, we propose a new direction for handling data flows on the network, which draws upon the knowledge – both theoretical and practical – we have acquired on the Internet, and borrows at the same time from fields such as databases. We propose some foundations for a decentralized architecture, allowing independent actors to exhange data and form an ecosystem built around the exchange of value. Although an ecosystem of potential entities already exists today, it is fair to say that no infrastructure allows its materialization yet.

This is the intent of this paper to propose an infrastructure allowing to put data as a first-class citizen of networks. We name it *Data-Aware Networks*, in the sense that the network will be able to route the queries to the appropriate hosts able to answer it, and support the transit of the data back to its requester. Just like packets are the basic unit considered on IP networks, grouped consistently within flows, we propose to consider the transport of data units triggered by a query, eventually involving the cooperation of several independent, federated entities.

The foundations of such Data-Aware Networks are presented in Section **??**, and the requirements for a communication protocol allowing the transport of queries between the different entities is proposed in Section **??**. We then consider how queries are handled within a single domain in Section 4, and the routing mechanisms we propose in Section 5. We then illustrate our proposal through some sample applications, and present MANIFOLD, an implementation of the concept that is powering two applications, namely TopHat [**?**] and MySlice [**?**]. Finally, Section **??** presents some related work. **Shall we announce the advantage of such an architecture already here ?**

2. FOUNDATIONS OF DATA-AWARE NETWORKING

2.1 The ecosystem of data

When we said previously that there is not a materialized ecosystem for data exchange in the network, it is more correct to notice that there is in fact a multitude of sparse efforts, each focused on a particular application. Let's look at Internet measurements, as an illustration. So far, there exist many independent efforts to measure its many facets, and more and more the need arise to be able to get access to all those different sources of data and combine them to extend our comprehension of the way it works. For various reasons, it would be impossible for a single entity to have the necessary knowledge, or to afford to build a global and versatily measurement platform. Instead, we see many efforts to propose datasets and real time data streams, in various formats, that can be used as building blocks for creating more advanced systems.

The collaboration between those entities relies mainly on their mutual interest, which mean that both find some value in mutualizing those resources (complementarity, scalability, sustainability, scale economies, etc.). **We could refer to a paper about federation...** Though, this collaboration is hindered by a vast set of technical barriers preventing an easy interconnection: data is often not consistently stored or retrieved, it is named differently, lacks some units, etc. Successful integration thus relies on a large amount of technical work, which is also what prevents scientific datasets to be easily and openly shared. We believe that those tasks, while not being fully transparent, should benefit from an agreed infrastructure and some support from the network.

It is of course a pipe dream to expect that it is possible that the network implement specific protocols for each type of data to be transported, or every way to contribute value to the ecosystem. We can think of entities providing measurements, caching, large storage, computing resources, aggregation and translation services, etc. But there are thousands more, and nothing should prevent a new entity to emerge if tomorrow it is able to bring some value for at least one member of the community[1]. We could make the analogy with the IP networks whose openness have allowed Content Providers to emerge and find they way in the ecosystem.

It seems natural to consider to what extent such an architecture can be extended and adapted to be applied to data, since its prevalence today deserves some additional support. We are thus considering a fully decentralized architecture, relying on a set of independent domains contributing value, which are free to negotiate and set up peering agreement

---

[1]This could be materialized by an exchange of money for example

between themselves.

2.2 The status today

In order to understand the choices that lays the grounds for our work, we need to have a look on the status of data today, and the fundamental requirement it has to support. Let's start with a few observations.

Data is highly heterogeneous.

- databases, flat files, XML, webservices,
- many concurrent protocols and data formats
- units, sampling,
- temporal aspects, annotations
- data flows: archive vs streaming vs on demand measurements

Data is scattered across numerous entities.

- aggregation that make some ground for routing
- ecosystem will also naturally evolve (economic incentive).

Many unsuspected data are complementary.

The ecosystem is huge and highly dynamic.

2.3 Overview of Data-Aware Networks

The architecture we propose aims at tackling those different challenges. First, we propose and motivate the need for a standard way to represent and transport data between end hosts, able to cope with a vast diversity of applications. Then, because for various reasons, the data will be split and distributed across several domains, and several hosts within a domain, we propose an adapted query mechanisms, only relying on a high level description of the data. **This place the data at the center of the network, not its location or its ownership.** A routing layer will be in charge of mapping such queries with the different entities able to answer it, both within a single domain, or across the federation. We note though that it is possible to explicitly target or exclude some known platforms during the query. **Routing = added value.**
Minimal amount of information to be shared to allow for a global federation
We will insert a schema illustrating the different domains: intra- and inter-domain, policies and agreements, specialized link or common substrate
Here we address dynamic and scalability concerns. size of routing tables parallel overlays = interests/community only information at the edge of the network (see later redundant section).

How to we adress multiple hops ?

Economical incentive, model, value sharing, etc discussed in Section 9.

2.4 Representation of data

Relation model.

We adopt the format of the relational model, which represent data as a multidimensional tuple associating a value to a set of attributes. We group the set of tuples having the same attributes within a class, which is identified by a key (a subset of attributes uniquely identifying the tuple within the class). The attributes can be used to reference an other class instance through their primary key as usual. The major difference with the relational model is that we allows those objects to possess methods. Such a representation has proved to be general enough to accommodate the different types of data that will be represented.

A particularity of our model is that the objects are not defined by a single platform, but instead are created though the joint declarations of the various platforms. Thus the different properties of an object might be spread over several platforms. Also, an object can be hosted on several platforms at the same time (example of caches), through for the sake of simplicity, we will assume here that there are no contradictions.

Use of ontologies.

The framework we are proposing in this article is agnostic of the different information it transports. Thus since we want to be able to unify and join information originating from various sources, we somehow need the different entities to use the same naming scheme for similar objects. We decided to rely on ontologies, which formally describe some agreed and simplified representation of a domain of knowledge. For example, we present some efforts in expressing ontologies for measurements in Section **??**.

Such ontologies fit our need as they define a semantic graph involving objects together with their attributes, their relations and their inheritance properties. For the sake of simplicity, we will limit ourselves to the case where all exchanged information belongs to a single domain, and that all platforms agree on the same ontology. In practise, we might extend our model to support multiple ontologies, and we could for instance consider new actors in charge of translating or bridging the information originating for several domains.

In practice, such graphs are huge and might hold several thousands of concepts and attributes. It is both unrealistic to assume that all platforms will expose the full set of attributes, and would pose serious scalability concerns for the elements of our network if they were to be aware of the full graph. The platform will typically only expose a small fraction of these objects and attributes (like a view in a database). Assuming the platforms agree on the underlying semantic, we will see that this information will be sufficient so that the intermediate nodes are able to reconstruct the subset of interest from the semantic (see Section 5.

3. COMMUNICATION PROTOCOL

3.1 Addressing data

From a user point of view, the federation should be transparent and appear as a single domain. That means someone issueing a query should not be required to know which domain is responsible for it (the partitioning of information). Note that this does not prevent the annotation of data with its provenance for transparency issues.

Reasons: we do not know who holds the data, in which

table/partition; furthermore it can be duplicated with two different schemas.

It is not realistic to assume that we can assign addresses or identifiers to each data unit, and effectively use them for routing matching the user query with the related tuples. One major drawback is the lack of possible way of aggregating addresses. Instead, we will rely on addressing the data by their class, together with the set of attributes of interest. Providing we have the right description of platforms, this information will be sufficient to deduce the set of tables that need to be combined to reconstitue partitions and views.

At this step, we can make an analogy with data warehousing solutions, that typically allow users to retrieve data though a set of dimensions that enhance the data available for a given fact (the class in our case). Those attribute can eventually be hierarchically organized. This combination of fact and dimensions is often refered to as a star, or snowflake schema.

3.2 Querying the federation

The interaction with the different platforms will be performed by sending a *query* that will be answered by a *result*. We do not limit ourselves to queries retrieving data, but we consider a general Data Manipulation Language allowing to read or alter the set of data units located at each platform, as well as executing methods on them. We support the five basic operations: `Get`, `Create`, `Update`, `Delete`[2] and `Execute`.

The results associated to a query will consist in the set of tuples corresponding to the object matched by the query. A set of fields will identify the attributes that need to be returned to the user, and thus will characterize which tables have to be joined on the platform. It will be further possible to restrict the range of a query thanks to a filter on the value of certain fields (a predicate on the form: attribute, operator, value), that might in addition put some partitions aside. Finally, update queries will hold a set of attributes whose value have to be updated.

It will be necessary to contact all the tables whose attributes are refered here, and thus they will all be required for establishing the query plane. **R/W is also necessary for routing both on table and on attributes**

A temporal parameter is added to the query to allow for the repatriate of existing results at a given epoch, but also to trigger on-demand measurements or even schedule new ones in the future.

Data address = Class (R/RW) + set of fields (R/RW)

Query = Action, filters, [params,] fields

Selected objects can be formulated as : Action(Processing(Restrict(fields)))

Can be expressed as operations(join(union)) provided we are in normal form. Though we talk about normal form later.

Other operators are considered such as SORT, LIMIT and OFFSET that will make a consistent navigation possible in the collection, as well as offer the base for more efficient processing. Such operators as well as other extensions will be proposed in future work, and won't be further described in this document.

---

[2]often denoted CRUD, which stands for Create/Read/Update/Delete, or IN-SERT/SELECT/UPDATE/DELETE in SQL, or POST/GET/PUT-PATCH/DELETE in REST.

3.3 Communication between peers

Let's now consider two hosts, identified by their IP address, one issueing a query for the second one. We can consider two communication schemes, that we will denote as *pull* and *pull*. They form the base of our communication architecture, that will allow the implementation of several behaviours, as we will see in Section 6.

Pull scheme.

In the *pull* scheme, the answer to a query will be addressed back to the emitter through in the reverse direction of the query. That means the data will flow backwards in the query plane (a tree) up to its root, which is connected to the previous hop. This is either the emitter of the query in the single domain case, or a previous domain which routed the query until this point. That means that a router will have to maintain an active query table, so that it can map incoming answers to their source peer. **Analogy with flow tables in flow-aware routers, cf FAN, Caspian, Anagran.**

This behaviour is adapted for one-time queries, and follows the expected behaviour of the answer flowing back to the source in the reverse direction.

Push scheme.

Sometimes, we might want to consider asynchronous schemes. This is useful for supporting alerts, or callback, asynchronous or periodic queries. It can also be used to address the result to a third party.

This is also convenient when the routing is just forwarding the data several times without modifying it. The source then has better directly addressing the data to the original host to avoid wasting network resources.

In practice, data might still have to be pushed to one or many intermediary hops since they might be involved in some processing. This behaviour will be resolved during the query plane establishment. We note it will also require some information about the support of such functionality in the different platforms. **Do we want to push further the IP analogy and present a possible data format for a communication protocol built on top of IP ?**

4. INTRADOMAIN ROUTING: THE QUERY PLANE SEEN AS A ROUTE

In this section, we suppose we have a single entity providing data. For the sake of convenience and analogy, we will assume it consists in a SQL database, and we will voluntary restrict ourselves to retrieving data. We also suppose the range of possible queries is known by the user; it corresponds to the set of possible views over the database schema. **We need to define which views we are considering : LEFT JOINS.**

Answering a query will be done in two distinct phases: **query routing** and **query optimization**, that respectively consist in choosing a set of candidate routes, or query planes, and selecting the cheapest one.

4.1 Query routing

The major difference between our query format and SQL is that we do not specify in which table the different attributes can be found. Instead, we rely on our unique naming assumption, that will allow us to find the table that need

to be joined to answer the query. We will denote this step as **query routing**, and we will also consider the selection of necessary partitions.

Determine intermediary tables.

Given a set of requested fields, we will have: Route(fields) = JOIN(UNION(Tables))

4.2 Query optimization

Overview.

This specification in fact regroups a range of various query planes, whose efficiency will depend on the join order for example. Further more, when applying the different operators (selection, projection, etc) for processing the data, many query planes with various costs will be generated, and it will be the role of the **query optimization** engine to choose the cheapest one.

Cost model.

The notion of cost is voluntarily vague since it will depend on the user choice and the range of available information: order of operation and their cost (memory, cpu, storage, network), number of tuples to be materialized, etc. We refer the reader to the relevant litterature from the database community for an overview of the different techniques that are commonly used.

Simple heuristics.

Join order : left tree, and not bushy (no parallelism considered at the moment). Some joins need to be done before others. eventually, the one involving as few tuples as possible ?

Then we will consider a series of equivalent query plans applying algebraic rules. Heuristic that consist in replacing the current query plan with an equivalent one with a cost supposed inferior.

4.3 Dataflows

Database network convergence. in the context of adaptive dataflows. query processing = adaptive dataflow

[Hellerstein: Cosmic convergence] implementation : pipelining operators. subexpression materialized. implementation more flexible than algebra. impose a router.

query plane will consist in a data flow.

4.4 Database normalization

In this section we describe the interest of database normalization for queries. **Later in routing, we will discuss about the interest of the normal form for routing : fewer changes, formula, etc.**

5. ROUTING DATA QUERIES

5.1 Peering

So far, we presented how a single entity can answer a query. In order to fully exploit the value offered by the federation, it is desirable that a platform is able to relay a query to another one, or even relying on the combined information from several platforms to answer its query. Like for the interdomain agreements on the internet, we will suppose that the different domains will negociate peering agreements, allowing them to interconnect and establish sessions with other peers.

The different domains will remain autonomous in the sense that they can use the technology of their choice for the routing of queries within a domain (cf the example of the SQL database in the previous section); they just need to expose standard interfaces to the federation, like it is the case in the internet with the BGP protocol.

In this section, we complement the communication and transport protocol we have presented so far by an announcement protocol, and a routing process. The former is used by the different platforms to disseminate information about the routes toward the data, in other terms the subset of the global schema that is accessible through this peer. The latter is used to efficiently answer the incoming query by choosing the best route, in term of a cost metric that is to be defined.

We expect this proved scheme to be adapted to our current purposes, since it allows the different platforms to keep their entire autonomy. Like in the internet, they are free to implement outbound and inbound policies that will select which information they are keen on sharing and receiving from the peers.

SCHEMA HERE
Do we need to say more about policies ?

5.2 Announce

Overview.

What can be put into queries, the routing will make the necessary glu.

Announce protocol.

Route plus cost.

Announce protocol

Quick reference to policies

Normalization.

5.3 Routing

Distributed query plane.

Convergence ? Optimality ?

Routed to each data unit, query plan, like multicast group. = Intradomain routing.

Bring the right query to the domain responsible to answering it.

We need a schema with the overview

analogy with a multicast group

like establishing a circuit.

5.4 Scalability

Only the useful part of the schema, only by the hosts that have to route and transport it. No initial knowledge of the schema, which can evolve dynamically according to the need. Routing only occurs at the edge of the network, done by the different domains. We envisage an eventual support from the network in Section 9.

dynamic aspects, reoptimization

6. APPLICATIONS

In this section, we illustrate some possible utilization of our framework through a set of representative use cases that we can find into other measurement federation proposals.

They are far from being exhaustive, and we in fact expect that this architecture allows for the emergence of unforeseen new actors bringing value into the system.

Distributed applications.

partitioning: space / time share responsibility for storing different classes of measurements eventually different domains : one way of splitting among many others

collaborative monitoring

Reactive measurements.

alarm on demand DSL : trigger -> callback

communication layer between services

storage service.

Get: traceroute, callback to storage on data, create !

or find a storage place

aggregation of measurement data: measure various complementary facets of the Internet, complementary, need for reactive meausrmeent to trigger measurements form a signal

contextualization of measurement data: in testbeds, measurements about testbed resources or its substrate (eg. planetlab nodes and public internet) or user experiment: requires some deductions

diagnostic of networks (perfSONAR): end to end path, require union of various results.

real time / streaming troubleshooting, monitoring : alarms, events. heterogeneous, collaborative and automatic allow rules and trigger to describe relationships

architecture allowing to distribute horizontally and vertically (partition vs functionalities):

7.   IMPLEMENTATION AND USE CASES

Could be implemented as a protocol on top of IP, or as a new network paradigm. Though for our current purpose, we implemented it as an overlay on top of the current network.

7.1   The MANIFOLD component

Implemented as an overlay to the current internet.

7.1.1   API

7.1.2   GUI

7.1.3   Authentication

7.2   OneLab services: TopHat & MySlice

8.   RELATED WORK

Ontologies

Example of measurement ontologies.

perfsonar measurement ontologies (openlab, novi, etc.) IPPM

Architecture

perfsonar SOA "fixed" set of actors, even though some are generic standardize both content and implementation format complex protocols for various tasks no routing, set of lookup services user acccess everything: based on identity federation only consider partitioning of data (UNION): no joining all the complexity exposed to user. different policy architecture

SFA OMF FRCP

no formal base support for network ???? nothing to standardize about dynamic aspects of the data and actors. Nothing has to be known by the network !!!!

Semantic = end hosts,

Allow many perspectives, some are described here

9.   EXTENSIONS TO THIS WORK

Formalization of the problem

Flow/Congestion control

packets protocols just like in the standard internet Could even be implemented on top of IP

Support from the network

routing at the edge, processing could be done optimally in the network benefit from many users, etc duplication query plane optimization etc.

Temporal aspects

Temporal consistency cache on-demand measurements Future measurement = schedule

Transparency and provenance

Security

Authentication, authorization

Economical models

What will the ecosystem look like ? On more generic objects ? how to share revenue cf panos (diversity) cf internet, AS relationships

impact on the shape (hierarchical, flat, etc.) we do not know what are incentives for people, what can have value for them.

APPENDIX

A. TOPHAT & MYSLICE

UPMC is running the TopHat and MySlice services, targeting two community of users, respectively the measurement and the experimenter ones. Both rely on the MANIFOLD interconnection framework that will be described in Section **??**.

A.1 TopHat

Overview

The TopHat measurement service has the dual role of collecting and aggregating measurements related to the PlanetLab testbed overlay to serve both the measuremnt community, but also supports experimenters from the testbed with measurements. It offers a dedicated service that provides network topology information to PlanetLab users and applications running on PlanetLab.

TopHat supplies measurements to its users thanks to its own TDMI platform, which are supplemented by drawing upon several independent specialized measurement and monitoring infrastructures, some of which have a proven track record of excellence in providing specialized measurements to the research community.

TopHat's niche lies in its ability to support of the entire lifecycle of an experiment: assisting users during experiment setup in choosing the nodes on which the experiment will be deployed; at run time, it providing live information to support adaptive applications and experiment control; and supporting retrospective analysis through access to archived data.

The tool architecture (see Section **??** allows it to support in the same way the ongoing testbed federation efforts, and allow measurement and monitoring integration across the different domains. It provides both an archive of regular measurements of the topology (the testbed substrate) and of the testbed infrastructure itself.

TDMI

The TopHat Dedicated Measurement Infrastructure (TDMI) is TopHat's own measurement infrastructure. It consists in modular probing agents that are deployed in a slice of the various PlanetLab nodes and probe the underlying network in a distributed efficient manner. In addition, they probe outwards to a number of target IP addresses that are not within PlanetLab. The aim of TDMI is to obtain the basic information required by TopHat. It implements such algorithms as Paris Traceroute to remove the artifacts arising from the presence of load balancers in the Internet. TDMI aims at providing the necessary information to TopHat users about the evolution of the overlay, and focuses on catching the dynamic aspects of the topology.

More details about the available data are presented in Section **??**.

Interconnected platforms

TopHat currently provides access to a set of interconnected system of different types:

measurement systems : TMDI, Gulliver, Dimes, Etomic, SONoMA

system-level monitoring infrastructure and testbed resources : CoMon, CoTop, MyPLC, Monitor, SFA

misc. information sources for topological and geographical data : Team Cymru IP to ASN mapping; MaxMind Geolite City; GeorgiaTech AS taxonomy

As TopHat relies on the modular and extensible MANIFOLD framework, it is possible to develop new platforms adapters to extend its reach.

A.2 MySlice

MySlice [MYSLICE1, MYSLICE2, MYSLICE3] consists in a user centric tool to support user's interaction with testbeds. It is tailored to support the full experiment lifecycle, from setup through completion, and it also includes the support for the management of a local authority (register and manage user accounts). Like TopHat, it is based on an open and extensible framework, as such it represents a candidate to provide a portal to the testbed federation.

The challenge MySlice solves is that the different platforms of interest for an experimenter typically realize one given functionality (SFA deals with the control plane, OMF with the experimental plane, TopHat with measurements, etc.) which is decoupled from users' needs at the various stages of their experiment. More specifically, as measurements as concerned, MySlice's objective is to put measurements into a context (measurements about a testbeds, a resource, a slice, etc.), and this is made possible thanks to the unifying framework.

MySlice currently provides the experimenter with measurements related to resources while he is selecting resources to be included into the slice, and the possibility to request new on demand measurements while the experiment is running. Finally, the possibility to collect and visualize measurements related to a slice is under development. As MySlice in itself is not the purpose of this document, the reader will refer to the referenced publications for further information about the tool.

A.3 Commonalities

• leverage a large ecosystem of available complementary and overlapping services and tools (far beyond testbed borders) • from our experience the UI is essential to users: need provide a transparent and consistent access • Exploit commonalities in platforms and processes

balance uniformity and heterogeneity
see paper FOSTER
rely on same common abstraction; MANIFOLD !

A.4 A convenience service

= a router at the edge Convenience service, deployed for a community not mandatory, only a service that can be short-circuited rely both on the MANIFOLD implementation we describe in the following, tailored and customized via the gateways (data and functionality) and plugins (web interface) same functionality works on user side also partially centralized service might bring additional benefits (value, latest version, etc), negociate access to sources

B. MANIFOLD: ARCHITECTURE AND COMMUNICATION PROTOCOL

B.1 Objects and collections

In MANIFOLD, we adopted an object oriented model, which has proved general enough to accommodate the need of

the different interconnected platforms so far. A user will typically handle a **collection**, a set of objects. An **object** consists of a *class* (a type) (like slice, resource, traceroute), etc.), one or more *fields*, among which a *key* that will uniquely identify the object instance, and zero or more *methods*. A property will eventually consists in a sub-collection (a slice hold a set of associated resources, a traceroute holds a set of hops, etc.).

A **field** will consist of a name, a type/class, a flag to determine whether it is an array, a read-only or read-write flag, and a description.

B.2 Building collection with relational operators

We voluntarily limit queries to a single class of objects.

The *timestamp* and *callback* elements are optional and respectively allow to specify a time range for the query (timstamp, interval, keywords), as well as a return channel (for long, asynchronous, periodic or large queries, or to be notified of some events). They will be detailed in further documents.

A query might also need to transport authentication token. This will be detailed in Section D.3.

B.3 Platform adapters: the notion of gateways

Our framework assumes that all platform expose a consistent query interface, and result results in the expected format. Unless some of them natively conform to this specification, we propose a mechanism of adaptors, named **gateways**.

A gateway will translate queries in the format that is expected by the platform, and map results into collections. Gateways will thus be in charge of handling issues related to transport and data representation. In addition, it might also adapt the semantic used by the platform when it differs from the adopted one.

We call *metadata* the minimal set of information that has to be provided by the platform or the gateway to allow for federation. This of course consists in the object description, which we extend with flag informing whether the platform supports fields, filters or any other operator. This will be of particular interest in the mediator when we have to determine which query can be sent to the platform, and which treatments are to be performed locally afterwards.

In some cases (a SQL database interface) it is possible to determine metadata automatically from the platform: the geteway will this propose such a functionality.

C. INTELLIGENT MEDIATOR

In this section we illustrate how the different components work together to provide access to the different sources of data. For simplicity, we consider four platforms (A to D) which all offer a single method. A method provides access to a set of measurements, each consisting in a set of fields identified by a unique colour. This representation is somehow similar to a relational database, and allows modeling for a wide range of sources of information, from web services to flat files.
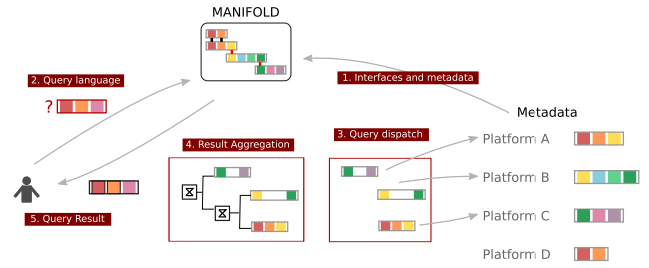

Figure 1: Illustration of a MANIFOLD query

1. Interfaces and metadata.

Each platform is described in a metadata files which is advertised to TopHat. This XML document gives general information about the platform, the way to connect to it, and the different methods and fields offered. Some fields uniquely identify a measurement and can be used for merging the data originating from several method calls, eventually from different platforms. This allows TopHat to form a global schema of available information. We note that this step requires a proper characterization of the ontologies used by the different platforms.

2. Querying TopHat.

TopHat proposes a simple query language to access the aggregated information. A query consists in a method, a temporal information, a filter pattern and a set of fields of interest, and we notice again the similarity with relational approaches. The methods, filters and fields are dependent on the interconnected platform and their availability, and it is important to notice that this range of these parameters is thus loosely defined, even though it might be possible to restrict it to a given ontology in future work.

3. Query dispatch....

TopHat will issue a set of queries to the different interconnected platforms to answer the user query in an optimal way. Since there might be redundant information (we assume no conflict), the platform will automatically choose the best combination of subqueries, but will also propose ways to guide its decisions.

3. ... and aggregation.

The result of the different subqueries will be joined together (eventually resolving time consistency issues thanks to the timing specifier of the user query), and send the result back to the user in the expected format. In addition to the fields provided by the platforms, TopHat introduces a few more ones to get information about the actual routing and aggregation performed: it is thus possible to annotate the different pieces of information with their origin, and their corresponding timestamp.

[WHY THE INTELLIGENT MEDIATOR IS GOOD FOR USERS : NO WORRY ABOUT WHO IS PROVIDING THE DATA/SERVICE, JUST QUERY FOR IT] [USED TO IMPLEMENT API and especially the GUI, GOOD FOR PLUGINS] [EXPLAIN THE DATA FLOW]

D.   THE MANIFOLD IMPLEMENTATION

NOTE: there are many possible implementation of this framework

D.1   Overview

MANIFOLD provides an implementation of the different functionalities we have presented so far: standard interfaces for querying the system or individual platforms for results; a set of gateways handling various platforms; the mediator functionality that allows to dispatch and route a query to the set of appropriate platforms, extending the value of the individual platforms; support for various authentication methods; a modular and extensible interface for the gui, offering visualization of results, and allowing to balance the diversity of results we get (the real value of the federation) and a uniform-enough presentation to the user.

D.2   Interfaces

MANIFOLD proposes a wide-range of interfaces to accommodate the diversity of users' needs:

a core library (written in Python) brings most of the MANIFOLD functionality, and exposes the query interface we have presented previously. Such a library can be conveniently integrated into third party projects in order to benefit from **MANIFOLD** the advanced mediator functionality, or simply from existing adapter written for it;

an API layer (currently XMLRPC, JSON and XMPP) exposes the core functionalities remotely via the same interface to which an authentication token is added (see below);

a web frontend : for most users needs, especially new users, a graphical user frontend will be the most convenient choice. It will allow for extended interaction and visualization with the data originating from the federation. MANIFOLD currently proposes a Joomla-based PHP frontend, and a Django version written in Python is being alternatively developed. Both rely on javascript over the jQuery framework for handling dynamic content. They integrate a plugin interface that allows for the customization of the interface, and more convenient way to display platform-specific information.

D.3   Authentication

In general, the different platforms interconnected will all have different authentication schemes, and thus will require different tokens. Myslice supports multiple users and allows to store for each of them the set of necessary token to contact the different platforms.

Despite federation efforts, it is likely that support for multiple authorization schemes will remain necessary. Consider as an example the PlanetLab Europe testbed, deployed over the public internet; the user might be interested in measurements or other kind of information provided by sources run by third parties well beyond the testbed borders, and the probability that all interconnected platforms support the same authentication scheme is close to zero.

Myslice is designed to support multiple users. It can be used either locally (python library) or remotely (API or GUI). In a local setup, a user does not need to authenticate to Myslice, while this is made mandatory for remote access by the API.

The following figure succintly presents how platforms, users and accounts are handled by MANIFOLD in its internal database:


Figure 2: MySlice database

Authentication to Myslice.

This functionality is used when users are remotely accessing MANIFOLD, either through the web GUI or the API.

Currently three authentication methods are supported:

- users have a local account

- users enter their OneLab/PlanetLab Europe or PlanetLab Central authentication token

- users enter their signed SFA certificate in their browser (only supported by the web GUI at the moment)

This scheme could be extended to more advanced solutions in the future, should the need arise (eg. SSO, etc.).

Authentication and authorization to remote platforms.

For each user, a set of accounts store the necessary information to authenticate and/or get authorized on the different platforms.

Supported methods include:

- anynymous access;

- unique authentication token (handled by the system, eg a login/password);

- user authentication token (better, since it allows for accounting);

- user delegated credentials (based on SFA, priviledged mechanism allowing accounting and higher trust).

interface between user tools and measurement and heterogeneous monitoring sources tight integration into MySlice protocol and standard representation format: query/results rely on a proper semantic definition framework to propose platform adapters

D.4 GUI architecture

TODO cf Loic's section on the TRAC

Platform for development: convenient aggregation and interoperability layer between the various services and the UI - async query mechanism - transparent access to data and functions - a uthentication information - caching and query optimization

functionality/data not present/etc.

E. USAGE MONITORING IN A FEDERATED ENVIRONMENT

- reporting EU, sponsors, public

- utilization of the platform

- its value in a federation [CONEXT]

- health / debug / FLS

- long term engineering tasks

- support experiments, research about measurements itself

- events, reactive measurement, probing

F. RELATED WORK

SFA = distributed architecture for secure transmission of assertions between recognized members.

Reuse as much as possible from the solutions that have emerged from the testbed federation problem. eg. SFA and OMF, lesson learned = separated process from semantic (testbed details). Let's show everything is always the same..

In fact Fits between a broader scheme of supporting a federation.

cf FRPC

Here a very lose scheme where we are not making any assumption We hope we can define a semantic. no notion of lifecycle. What does it mean for measurement ?

Why keeping the SFA then ?

perfsonar: cf perfsonar notes, multiple comparisons here

[TOPHAT] [MYSLICE-FOSTER] [MYSLICE-GEC] [MYSLICE-FEW]

G. REFERENCES